

Task 1. Data Preparation for Modelling

1. What variables did you include in the modelling, and what were their roles and measurement level set? Justify your choice.

Variable	Role	Measurement level	Reason of usage
Auction	Input	Categorical	<ul style="list-style-type: none"> Transaction standards may differ for each auction company. There is a possibility that the defect rate is high in certain auctions.
VehYear	Input	Numerical	<ul style="list-style-type: none"> Vehicle year reflects the car's age and condition Older vehicles are more likely to have hidden issues leading to a kick.
Transmission	Input	Categorical	<ul style="list-style-type: none"> Transmission type (automatic, manual) affects demand and price differences, which can correlate with bad buy risks. Most consumers have a tendency to use an automatic type of car.
VehOdo	Input	Numerical	<ul style="list-style-type: none"> A vehicle range finder indicates how much the vehicle has been used The higher the mileage, the greater the risk of mechanical failure.
Nationality	Input	Categorical	<ul style="list-style-type: none"> Country of manufacture can imply differences in quality, reliability, and market perception, impacting kick probability.
IsBadBuy	Target	Numerical	<ul style="list-style-type: none"> The outcome variable to be predicted: 1 = bad buy (kick), 0 = not a bad buy
MMRCurrentAuctionAveragePrice	Input	Numerical	<ul style="list-style-type: none"> Auction market price in average condition Large discrepancies with the actual price may indicate abnormal transactions.
MMRCurrentRetailAveragePrice	Input	Numerical	<ul style="list-style-type: none"> Retail market price in average condition If too close to or lower than auction price, it signals potential risk of a kick.

2. Did you have to fix any data quality problems, including data imputation? Detail them.

Auction		Transmission	
<pre>===== Column: Auction ===== Before preprocessing: Auction MANHEIM 22168 ADESA 11086 OTHER 8178 NaN 44 Name: count, dtype: int64 ===== After preprocessing: Auction MANHEIM 22168 ADESA 11086 OTHER 8222 Name: count, dtype: int64</pre>	<ul style="list-style-type: none"> NaN (empty value) should be changed to 'OTHER' value. Code: <pre>df['Auction'] = df['Auction'].fillna('OTHER')</pre> 	<pre>===== Column: Transmission= Before preprocessing: Transmission AUTO 39930 MANUAL 1495 NaN 44 ? Manual 1 Name: count, dtype: int64 ===== After preprocessing: Transmission AUTO 39930 MANUAL 1496 UNKNOWN 50 Name: count, dtype: int64</pre>	<ul style="list-style-type: none"> String was changed to upper characters. '?' value → NaN → 'UNKNOWN' Code: <pre>df['Transmission'] = df['Transmission'].str.upper() df['Transmission'] = df['Transmission'].replace('?', np.nan) df['Transmission'] = df['Transmission'].fillna('UNKNOWN')</pre>
Make		Nationality	
<pre>===== Column: Make ===== Before preprocessing: MERCURY CHEVROLET 9548 DODGE 7385 FORD 6943 CHRYSLER 5259 PODUSLIC 2293 KIA 1337 SATURN 1248 NISSAN 1186 DEZER 995 HYUNDAI 957 SUZUKI 842 TOYOTA 664 MITSUBISHI 569 MAZDA 533 HONDA 527 BUICK 432 GMPC 351 NINJA 284 OLDSMOBILE 146 TESLA 93 SCION 77 VOLKSWAGEN 73 LINCOLN 58 NAVI 44 ZINCHINTZ 27 MINI 19 ACURA 18 CADILLAC 17 SUBARU 17 LEXUS 15 VOLVO 12 Name: count, dtype: int64</pre>	<ul style="list-style-type: none"> NaN → 'UNKNOWN' Code: <pre>df['Make'] = df['Make'].fillna('UNKNOWN')</pre> 	<pre>===== Column: Nationality= Before preprocessing: Nationality AMERICAN 34615 OTHER ASIAN 4474 TOP LINE ASIAN 2110 USA 125 OTHER 104 NaN 44 ? 3 Name: count, dtype: int64 ===== After preprocessing: Nationality AMERICAN 34741 OTHER ASIAN 4474 TOP LINE ASIAN 2110 OTHER 104 UNKNOWN 47 Name: count, dtype: int64</pre>	<ul style="list-style-type: none"> 'USA' → 'AMERICAN' ? → NaN → 'UNKNOWN' Code: <pre>df['Nationality'] = df['Nationality'].str.upper().str.strip() df['Nationality'] = df['Nationality'].replace('USA', 'AMERICAN') df['Nationality'] = df['Nationality'].replace('?', np.nan) df['Nationality'] = df['Nationality'].fillna('UNKNOWN')</pre>

TopThreeAmericanName	ForSale
===== Column: TopThreeAmer Before preprocessing: TopThreeAmericanName GM 14075 CHRYSLER 13627 FORD 7039 OTHER 6688 NaN 44 ? 3 Name: count, dtype: int64 ===== After preprocessing: TopThreeAmericanName GM 14075 CHRYSLER 13627 FORD 7039 OTHER 6688 UNKNOWN 47 Name: count, dtype: int64	===== Column: ForSale===== Before preprocessing: ForSale Yes 27402 YES 8544 yes 5524 ? 3 No 2 0 1 Name: count, dtype: int64 ===== After preprocessing: ForSale 0.0 41470 UNKNOWN 4 1.0 2 Name: count, dtype: int64
===== Column: IsBadBuy===== Before preprocessing: count 41476.000000 mean 0.129497 std 0.335753 min 0.000000 25% 0.000000 50% 0.000000 75% 0.000000 max 1.000000 Name: IsBadBuy, dtype: float64 ===== After preprocessing: count 41476.000000 mean 0.129497 std 0.335753 min 0.000000 25% 0.000000 50% 0.000000 75% 0.000000 max 1.000000 Name: IsBadBuy, dtype: float64	===== Column: VehYear===== Before preprocessing: count 41432.000000 mean 2095.360615 std 1.738587 min 2001.000000 25% 2094.000000 50% 2095.000000 75% 2097.000000 max 2019.000000 Name: VehYear, dtype: float64 ===== After preprocessing: count 41476.000000 mean 2095.360232 std 1.729769 min 2001.000000 25% 2094.000000 50% 2095.000000 75% 2097.000000 max 2019.000000 Name: VehYear, dtype: float64
MMR Price	VehYear
	<p>• This variable does not need imputation</p> <p>• Strip: remove 'blank space'</p> <p>• '?', 0 → NaN → median value</p> <p>• Change to number</p> <p>• Code:</p> <pre>mm_current_cols = ['MMRCurrentRetailAveragePrice', 'MMRCurrentAuctionAveragePrice', 'MMRCurrentRetailAvgPrice', 'MMRCurrentAuctionAvgPrice', 'MMRCurrentRetailClearPrice', 'MMRCurrentAuctionClearPrice', 'MMRCurrentRetailAvgClearPrice', 'MMRCurrentAuctionAvgClearPrice', 'MMRCurrentRetailClearAvgPrice'] for col in mm_current_cols: df[col] = df[col].apply(str).str.replace(' ', '') df[col] = df[col].replace('?', np.nan) df[col] = pd.to_numeric(df[col], errors='coerce') df[col] = df[col].replace(0, np.nan) df[col] = df[col].fillna(df[col].median())</pre>
According to the heat map, there are nine variable regarding MMRPrice, and these variables high correlation is confirmed in correlation matrix graph. Thus, we chose two variables: MMRCurrentRetailAveragePrice, MMRCurrentAuctionAveragePrice	<p>• If value is lower than lower_bound and is higher than upper_bound, these values replaced median values.</p> <p>• That is why outliers can affect incorrect analysis</p> <p>• Code:</p> <pre>lower_bound = 1888 upper_bound = 400000 vehodo_median = df['VehODO'].median() df.loc[(df['VehODO'] < lower_bound) (df['VehODO'] > upper_bound), 'VehODO'] = np.nan df['VehODO'] = df['VehODO'].fillna(vehodo_median)</pre>

3. Report the proportion of values of the target variable for the dataset before and after the pre-processing.

Nationality	TopThreeAmericanName	ForSale
<pre>*** Nationality *** *** Before preprocessing *** IsBadBuy Nationality 0 AMERICAN 0.837386 OTHER ASIAN 0.196857 TOP LINE ASIAN 0.050406 USA 0.002939 OTHER 0.002329)) 0.000883 1 AMERICAN 0.822748 OTHER ASIAN 0.115564 TOP LINE ASIAN 0.054427 OTHER 0.003728 USA 0.003541 Name: proportion, dtype: float64 ----- *** After preprocessing *** IsBadBuy Nationality 0 AMERICAN 0.839441 OTHER ASIAN 0.196744 TOP LINE ASIAN 0.050353 OTHER 0.002327 UNKNOWN 0.001136 1 AMERICAN 0.825358 OTHER ASIAN 0.115435 TOP LINE ASIAN 0.054566 OTHER 0.003724 UNKNOWN 0.001117 Name: proportion, dtype: float64 -----</pre>	<pre>*** TopThreeAmericanName *** *** Before preprocessing *** IsBadBuy TopThreeAmericanName 0 GM 0.346854 CHRYSLER 0.331356 FORD 0.162115 OTHER 0.159592 ?) 0.000883 1 CHRYSLER 0.312395 GM 0.291705 FORD 0.222181 OTHER 0.173719 Name: proportion, dtype: float64 ----- *** After preprocessing *** IsBadBuy TopThreeAmericanName 0 GM 0.346489 CHRYSLER 0.331077 FORD 0.162044 OTHER 0.159424 UNKNOWN 0.001136 1 CHRYSLER 0.312846 GM 0.291388 FORD 0.222193 OTHER 0.173524 UNKNOWN 0.001117 Name: proportion, dtype: float64 -----</pre>	<pre>*** ForSale *** *** Before preprocessing *** IsBadBuy ForSale 0 Yes 0.666833 YES 0.202714 yes 0.130287 ?) 0.000883 1 No 0.000055 0 0.000028 Yes 0.619252 YES 0.228077 yes 0.152672 Name: proportion, dtype: float64 ----- *** After preprocessing *** IsBadBuy ForSale 0 0.0 0.999834 UNKNOWN 0.000111 1.0 0.000055 1 0.0 1.000000 Name: proportion, dtype: float64</pre>
VehYear	VehOdo	MMRCurrentAuctionAveragePrice
<pre>*** VehYear *** *** Before preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36067.0 2085.471151 1.697094 2081.0 2080.0 2086.0 2087.0 1 5365.0 2084.617521 1.770015 2081.0 2083.0 2085.0 2086.0 max IsBadBuy 0 2018.0 1 2009.0 ----- *** After preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36105.0 2085.470655 1.696270 2081.0 2084.0 2086.0 2087.0 1 5371.0 2084.617948 1.769872 2081.0 2083.0 2085.0 2086.0 max IsBadBuy 0 2018.0 1 2009.0 -----</pre>	<pre>*** VehOdo *** *** Before preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36067.0 70848.963263 14737.768855 577.0 61073.0 72355.0 1 5365.0 74332.241193 14267.248465 4825.0 65222.0 76401.0 max IsBadBuy 0 81832.5 489444.0 1 84692.0 115717.0 ----- *** After preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36105.0 70842.090486 14566.926787 8706.0 61099.0 72655.0 1 5371.0 74330.896481 14259.332414 4825.0 65226.5 76390.0 max IsBadBuy 0 81820.0 112056.0 1 84687.5 115717.0 -----</pre>	<pre>*** MMRCurrentAuctionAveragePrice *** *** Before preprocessing *** count unique top freq IsBadBuy 0 36064 8795 0 251 1 5365 3647 0 36 ----- *** After preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36105.0 6250.856488 2167.090173 2017.0 4558.0 6138.0 7757.0 1 5371.0 5480.379352 2206.092813 2017.0 3648.5 5232.0 6951.5 max IsBadBuy 0 12627.0 1 12617.0 -----</pre>
MMRCurrentRetailAveragePrice		
<pre>*** MMRCurrentRetailAveragePrice *** *** Before preprocessing *** count unique top freq IsBadBuy 0 36047 18351 0 251 1 5362 3879 0 36 ----- *** After preprocessing *** count mean std min 25% 50% 75% \ IsBadBuy 0 36105.0 8920.540978 2762.134498 2973.0 6849.0 8760.0 10927.0 1 5371.0 7951.943400 2815.670232 2970.0 5760.0 7838.0 9839.0 max IsBadBuy 0 16151.0 1 16098.0 -----</pre>		<ul style="list-style-type: none"> After pre-processing, several variables showed skewed distributions in the Kick Car dataset, such as Auction MANHEIM (0.49), Transmission-AUTO(0.96), Nationality – American (0.83). The average values of VehYear, VehOdo, MMRCurrentAuctionAveragePrice, and MMRCurrentRetailAveragePrice tend to center around 50%

Task 2. Decision Tree Modelling

- Build a decision tree using the default setting. Examine the tree results and answer the following:
 - What parameters have been used to build the tree? Detail them.

Default tree (DecisionTreeClassifier)	Small model tree (DecisionTreeClassifier)
<pre>****Model parameters**** {'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': None, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'monotonic_cst': None, 'random_state': 10, 'splitter': 'best'} Number of leaves in the trained model: 4999</pre>	<pre>model_small = DecisionTreeClassifier(max_depth=3, class_weight='balanced', random_state=random_state)</pre>
'criterion': 'gini', min_samples_leaf: 1, min_samples_split: 2, random_state: 10	max_depth = 3, class_weight='balanced', random_state=random_state

b. What data split was used to create training and test datasets?

```
from sklearn.model_selection import train_test_split
random_state = 10
test_set_size = 0.3
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_set_size, stratify=y, random_state=random_state)

print("Size of training set:", len(X_train))
print("Size of testing set:", len(X_test))

Size of training set: 29033
Size of testing set: 12443
```

- 'random_state (10)' sets a fixed value to ensure that we get the same result regardless of multiple tries. Data for testing is divided by the ratio of 'test_set_size (0.3). Therefore, the size of the training set is 29033, and the testing set is 41476


c. What is the classification accuracy on the training and test datasets?

<pre>print("Training set accuracy:", model.score(X_train, y_train)) print("Testing set accuracy:", model.score(X_test, y_test)) Training set accuracy: 0.9998277821788999 Testing set accuracy: 0.7807602668166841</pre>	<p>Training accuracy shows high accuracy (0.99), whereas testing accuracy (0.78) is lower than training accuracy.</p> <p>This result shows the possibility of overfitting.</p>																														
<p>- default tree</p> <table style="margin-left: auto; margin-right: auto;"> <thead> <tr> <th></th> <th>precision</th> <th>recall</th> <th>f1-score</th> <th>support</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>0.88</td> <td>0.87</td> <td>0.87</td> <td>10832</td> </tr> <tr> <td>1</td> <td>0.18</td> <td>0.19</td> <td>0.18</td> <td>1611</td> </tr> <tr> <td>accuracy</td> <td></td> <td></td> <td>0.78</td> <td>12443</td> </tr> <tr> <td>macro avg</td> <td>0.53</td> <td>0.53</td> <td>0.53</td> <td>12443</td> </tr> <tr> <td>weighted avg</td> <td>0.79</td> <td>0.78</td> <td>0.78</td> <td>12443</td> </tr> </tbody> </table>		precision	recall	f1-score	support	0	0.88	0.87	0.87	10832	1	0.18	0.19	0.18	1611	accuracy			0.78	12443	macro avg	0.53	0.53	0.53	12443	weighted avg	0.79	0.78	0.78	12443	<p>Regarding 'class 1' from this classification_report,</p> <ul style="list-style-type: none"> - precision=0.18, recall=0.19, f1-score=0.18 <p>From this report, the decision tree model identified 19% of kick car (recall), and of those, 18% were correctly (precision). This means this model cannot predict class 1(kick).</p> <p>Thus, we effort to tune for improvement of model.</p>
	precision	recall	f1-score	support																											
0	0.88	0.87	0.87	10832																											
1	0.18	0.19	0.18	1611																											
accuracy			0.78	12443																											
macro avg	0.53	0.53	0.53	12443																											
weighted avg	0.79	0.78	0.78	12443																											
<p>- small model tree</p> <table style="margin-left: auto; margin-right: auto;"> <thead> <tr> <th></th> <th>precision</th> <th>recall</th> <th>f1-score</th> <th>support</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>0.91</td> <td>0.63</td> <td>0.74</td> <td>10832</td> </tr> <tr> <td>1</td> <td>0.19</td> <td>0.57</td> <td>0.28</td> <td>1611</td> </tr> <tr> <td>accuracy</td> <td></td> <td></td> <td>0.62</td> <td>12443</td> </tr> <tr> <td>macro avg</td> <td>0.55</td> <td>0.60</td> <td>0.51</td> <td>12443</td> </tr> <tr> <td>weighted avg</td> <td>0.82</td> <td>0.62</td> <td>0.68</td> <td>12443</td> </tr> </tbody> </table>		precision	recall	f1-score	support	0	0.91	0.63	0.74	10832	1	0.19	0.57	0.28	1611	accuracy			0.62	12443	macro avg	0.55	0.60	0.51	12443	weighted avg	0.82	0.62	0.68	12443	<p>The default tree has the complexity of model due to unlimited maximum depth, so this model can be enhanced by pruning branches of the decision tree (small model tree).</p> <p>We set up max_depth = 3.</p> <p>In the evaluation of model, the result shows that recall = 0.57, and f1-score = 0.28. However, precision remains at a low level.</p>
	precision	recall	f1-score	support																											
0	0.91	0.63	0.74	10832																											
1	0.19	0.57	0.28	1611																											
accuracy			0.62	12443																											
macro avg	0.55	0.60	0.51	12443																											
weighted avg	0.82	0.62	0.68	12443																											

d. What is the size of the tree (number of nodes and rules)?

<p>< default tree></p> <pre>print('Number of leaves in the trained model:', model.get_n_leaves()) print('Number of nodes:', model.tree_.node_count)</pre> <p>Number of leaves in the trained model: 4999 Number of nodes: 9997</p>	<p><small model tree></p> <pre># Node, Leaves, depth print("Nodes:", model_small.tree_.node_count) print("Leaves:", model_small.get_n_leaves()) print("Depth:", model_small.get_depth())</pre> <p>Nodes: 15 Leaves: 8 Depth: 3</p>
<p>Default tree (maximal tree) has 4999 rules and 9997 nodes, whereas small model tree has 8 rules and 15 nodes.</p>	

e. Which variable is used for the first split? What are the variables that are used for the second split?



f. What are the five important variables in building the tree?

< default model tree>	< small model tree>
<pre>display_feature_importances(model, feature_names) VehOdo : 0.3693214055782883 MMRCurrentRetailAveragePrice : 0.2659883719688045 MMRCurrentAuctionAveragePrice : 0.24943085804225154 VehYear : 0.042967744890411944 Auction_OTHER : 0.01542379634537651 Auction_MANHEIM : 0.0129631752570289 Auction_ADESA : 0.010884338151795542 Nationality_AMERICAN : 0.09273743181650274 Nationality_OTHER ASIAN : 0.009260489473204996 Nationality_TOP LINE ASIAN : 0.006488607453676483 Transmission_AUTO : 0.0036549128945540465 Transmission_MANUAL : 0.00329320799643804224 Nationality_OTHER : 0.001049866298554846 Transmission_UNKNOWN : 0.0 Nationality_UNKNOWNN : 0.0 Number of leaves: 4999</pre> <ul style="list-style-type: none"> - VehOdo: 0.36 - MMRCurrentRetailAveragePrice: 0.26 - MMRCurrentAuctionAveragePrice: 0.24 - VehYear: 0.04 - Auction_OTHER: 0.01 <pre>def display_feature_importances(model, feature_names, features_to_display=20): importances = model.feature_importances_ indices = np.argsort(importances) indices = np.flip(indices, axis=0) indices = indices[:features_to_display] for i in indices: print(feature_names[i], ':', importances[i]) print("Number of leaves:", model.get_n_leaves())</pre>	<pre>display_feature_importances(model_small, feature_names) visualize_model(model_small) VehYear : 0.76651573891371191 MMRCurrentAuctionAveragePrice : 0.11149221549044215 VehOdo : 0.0731100586607268 Auction_ADESA : 0.023442055720881594 Auction_MANHEIM : 0.01979828879091631 Nationality_TOP LINE ASIAN : 0.0 Nationality_UNKNOWNN : 0.0 Transmission_MANUAL : 0.0 Transmission_UNKNOWN : 0.0 Nationality_AMERICAN : 0.0 Nationality_OTHER : 0.0 Nationality_OTHER ASIAN : 0.0 Auction_OTHER : 0.0 Transmission_AUTO : 0.0 MMRCurrentRetailAveragePrice : 0.0 Number of leaves: 8</pre> <ul style="list-style-type: none"> - VehYear: 0.76 - MMRCurrentAuctionAveragePrice: 0.11 - VehOdo: 0.073 - Auction_ADESA: 0.029 - Auction_MANHEIM: 0.019
	<ul style="list-style-type: none"> ▪ There are five important variables according to default tree and small tree, this situation can be showed from <code>display_feature_importances</code>' function. ▪ Auction can be included in five important variables, but low percentage (0.01, 0.019) means that these data are not meaning variables.


g. Report if you see any evidence of model overfitting

Overfitting means it can fit the data in the training process regardless of the correct prediction result:

- **The number of nodes & complicated and deep model**
When confirming default maximal tree, there is a complicated and deep decision tree model, which has a lot of nodes.
- **Training accuracy (0.99) > testing accuracy (0.78)**
The maximal tree overfits the training dataset (0.99). That is why testing accuracy is lower than training. Therefore, this model fits in the case of training, not in the case of testing.

2. Build another decision tree tuned with GridSearchCV. Examine the tree results.

- a. What are the optimal parameters for this decision tree? Explain your choice of hyperparameters to search, and the chosen search range(s)?



b. What is the classification accuracy on the training and test datasets?

<pre>Fitting 10 folds for each of 20 candidates, totalling 200 fits Train accuracy: 0.6274584093962043 Test accuracy: 0.6234027163867235 precision recall f1-score support 0 0.91 0.63 0.74 10832 1 0.19 0.57 0.28 1611 accuracy 0.62 macro avg 0.55 0.60 0.51 12443 weighted avg 0.82 0.62 0.68 12443 {'criterion': 'entropy', 'max_depth': 3, 'min_samples_leaf': 5}</pre> <p><1> Overall, training (0.627) and test accuracy (0.623) are similar score. This result may mean overfitting risk is reduced, but it may also indicate underfitting.</p>	<p><2> Regarding ‘class 1’:</p> <ul style="list-style-type: none"> precision= 0.19, recall=0.57, f1-score=0.28 <p>From this report, the decision tree model identified 57% of kick car (recall), and of those, 19% were correctly (precision). This means this model cannot predict class 1(kick). F1-score (0.28) is a balance between precision and recall.</p> <p>Although recall can be improved, precision is still showed to low score, and f1-score also is low score due to this precision score</p>
--	---

c. What is the size of the chosen tree (number of nodes and rules)?

```
params = {
    'criterion': ['gini', 'entropy'],
    'max_depth': range(2, 4),
    'min_samples_leaf': range(5, 50, 10)
}
cv = perform_grid_search(X_train, y_train, X_test, y_test, params)

Fitting 10 folds for each of 20 candidates, totalling 200 fits
Train accuracy: 0.6274584099362043
Test accuracy: 0.6234027163867235
      precision    recall   f1-score   support
          0       0.91     0.63     0.74    10832
          1       0.19     0.57     0.28     1611

   accuracy        0.55     0.60     0.51    12443
  macro avg        0.82     0.62     0.68    12443
weighted avg      0.82     0.62     0.68    12443


{'criterion': 'entropy', 'max_depth': 3, 'min_samples_leaf': 5}
```

Nodes: 15
Leaves: 8
Depth: 3

We choose this gridserachCV model with these parameters; this model has 15 nodes and 8 rules, which is similar to a small decision classifier model (small model tree).

However, this model (entropy) has a difference in the criterion from the small model tree (gini).


d. Which variable is used for the first split? What are the variables that are used for the second split?



e. What are the five important variables in building the tree?

<pre>display_feature_importances(model_cv, feature_names) visualize_model(model_cv)</pre> <p>VehYear : 0.7598540857360376 MMRCurrentAuctionAveragePrice : 0.1144055334609266 VehOdo : 0.07466742514630023 Auction_ADESA : 0.03162454160340716 Auction_MANHEIM : 0.019448414053328502 Nationality_TOP LINE ASIAN : 0.0 Nationality_UNKNOWN : 0.0 Transmission_MANUAL : 0.0 Transmission_UNKNOWN : 0.0 Nationality_AMERICAN : 0.0 Nationality_OTHER : 0.0 Nationality_OTHER ASIAN : 0.0 Auction_OTHER : 0.0 Transmission_AUTO : 0.0 MMRCurrentRetailAveragePrice : 0.0 Number of leaves: 8</p>	<ul style="list-style-type: none"> - VehYear: 0.75 - MMRCurrentAuctionAveragePrice: 0.11 - VehOdo: 0.07 - Auction_ADESA: 0.03 - Auction_MANHEIM: 0.019
--	---

f. Report if you see any evidence of model overfitting.




3. What is the significant difference between these two decision tree models – default (Task 2.1) and using GridSearchCV (Task 2.2)? How do they compare performance-wise? Produce the ROC curve for both DTs. Explain why those changes may have happened

<default model-DecisionTreeClassifier>


- This model has 9997 nodes.
 - It is a complicated and deep model.

<Small model-DecisionTreeClassifier>



- gini = 0.499

<GridSearchCV model>



- entropy = 0.985

- Entropy is information gain.


- The higher the entropy, the much information (purity).

<ROC curve>

```
print("ROC index on test for default model:", roc_index_dt)
print("ROC index on test for small model:", roc_index_dt_small)
print("ROC index on test for grid search model:", roc_index_dt_cv)

ROC index on test for default model: 0.5288166393434355
ROC index on test for small model: 0.6386763716857975
ROC index on test for grid search model: 0.6386763716857975
```

- **Difference:** Small tree is based on 'gini', whereas cv model is based on 'entropy' (information gain)
 - **Comparison of performance:** The ROC curve can help compare models. These models' index is more than 0.5, and small tree and grid search is the same to 0.639. Therefore, we can choose a model, considering gini and entropy.
 - **Why these changes:** The differences between default tree and other models emerge because we set up 'class_weight=balanced' for reducing overfitting in small default model. Furthermore, in process of GridSerachCV tuning (cv.best_estimator_), model based on 'entropy' is considered better cross-validation quality.



4. From the better model, can you provide the characteristics of cars most likely to be ‘kicks’? If it is hard to comprehend, discuss why.

GridSearchCV decision tree (criterion: entropy) → can address important variables.

- Vehicle Year (First split: VehYear <=2005.5): cars that are older than 2005 have tendency to have a high risk of 'kick(badbuy)'
 - Vehicle Odometer (VehOdo <= 68528 and 75443.5): high odometer can be related to kicks
 - Auction (= MANHEIM, ADESA): Auction company influences the likelihood of kick.
 - MMRCurrentAuctionAveragePrice <= 6171: A low auction price can mean that this car has a high risk regarding kick (bad buy)

It is hard to comprehend because:

- The model is complicated, which has many leaves and rules.
- We cannot confirm the correlation with variables because many variables can simultaneously affect to model's results.
- Data imbalance: The data set is imbalanced with a significant number of negative data and a few positive data extremely.

Task 3. Regression Modelling

- Describe what additional processing was required on this dataset to be used in regression modelling. List the variables that need further processing and provide details of the processing.

• Standardisation


```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
print("Before scaling-----")
for i in range(5):
    col = X_train[:,i]
    print("Variable #{}: min {}, max {}, mean {:.2f} and std dev {:.2f}".format(i, min(col), max(col), np.mean(col), np.std(col)))
X_train = scaler.fit_transform(X_train, y_train)

print("After scaling-----")
for i in range(5):
    col = X_train[:,i]
    print("Variable #{}: min {}, max {}, mean {:.2f} and std dev {:.2f}".format(i, min(col), max(col), np.mean(col), np.std(col)))

X_test = scaler.transform(X_test)
Before scaling
-----
Variable #0: min 2001.0, max 2010.0, mean 2005.36 and std dev 1.73
Variable #1: min 4825.0, max 115717.0, mean 7131.23 and std dev 14570.21
Variable #2: min 2017.0, max 12623.0, mean 6156.95 and std dev 2188.21
Variable #3: min 0.0, max 1.0, mean 0.48 and std dev 0.27
Variable #4: min False, max True, mean 0.27 and std dev 0.44
After scaling
-----
Variable #0: min -2.5132851528044004, max 2.584588421782512, mean -0.08 and std dev 1.00
Variable #1: min -8.561035230327706, max 3.0476255036212663, mean -0.00 and std dev 1.00
Variable #2: min -1.891938523080326, max 2.954951367499558, mean 0.00 and std dev 1.00
Variable #3: min -2.0884404573778474, max 2.635112730794915, mean 0.00 and std dev 1.00
Variable #4: min -0.603342483025691, max 1.6574560499638853, mean 0.00 and std dev 1.00
```

- Numerical variables such as VehYear, VehOdo, MMRCurrentAuctionAveragePrice, and MMRCurrentRetailAveragePrice were standardised using StandardScaler.
- This process ensures mean = 0 and standard deviation = 1 and prevents variables with large-scale differences.

• Logarithmic transformation



- Variables such as VehYear, VehOdo, MMRCurrentAuctionAveragePrice, and MMRCurrentRetailAveragePrice show skewed distributions.
- These are normally distributed by applying a logarithmic transformation.

- `class_weight='balanced'`: Most of 'IsBadBuy' variable are 1 (x kick car) in 'kick.vsc' dataset. Therefore, it should be balanced with this processing.

2. Build a regression model tuned with GridSearchCV. Answer the following:

a. Name the Regression function used.

```
: from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state=random_state)

model.fit(X_train, y_train)
```

LogisticRegression  

LogisticRegression(random_state=10)

b. What are the optimal parameters for this regression model? Explain your choice of hyperparameters to search, and the chosen search range(s)?

<pre>from sklearn.model_selection import GridSearchCV params = {'C': [0.001, 0.01, 0.1, 1, 10, 100]} cv = GridSearchCV(param_grid=params, estimator=LogisticRegression(random_state=random_state, class_weight='balanced'), cv=10, n_jobs=-1) cv.fit(X_train, y_train) print("Train accuracy:", cv.score(X_train, y_train)) print("Test accuracy:", cv.score(X_test, y_test)) y_pred = cv.predict(X_test) print(classification_report(y_test, y_pred)) print(cv.best_params_) Train accuracy: 0.614404298556147 Test accuracy: 0.6149642269939568 precision recall f1-score support 0 0.91 0.62 0.74 10832 1 0.19 0.61 0.29 1611 accuracy macro avg 0.55 0.61 0.51 12443 weighted avg 0.82 0.61 0.68 12443 {'C': 1}</pre>	<pre>params = {'C': [0.001, 0.01, 0.1, 1, 10, 100]} cv = GridSearchCV(param_grid=params, estimator=LogisticRegression(random_state=random_state, class_weight='balanced'), cv=10, n_jobs=-1) cv.fit(X_train_log, y_train_log) print("Train accuracy:", cv.score(X_train_log, y_train_log)) print("Test accuracy:", cv.score(X_test_log, y_test_log)) y_pred = cv.predict(X_test_log) print(classification_report(y_test_log, y_pred)) print(cv.best_params_) Train accuracy: 0.618916405469638 Test accuracy: 0.620429150957181 precision recall f1-score support 0 0.91 0.62 0.74 10832 1 0.19 0.61 0.29 1611 accuracy macro avg 0.55 0.62 0.52 12443 weighted avg 0.82 0.62 0.68 12443 {'C': 0.01}</pre>
--	--

- Smaller C = Stronger regularisation, Larger C = Weaker regularisation → Ideal range = 10^{-6} and 10^4
- To be balanced between overfitting and underfitting, we chose search range for Hyperparameters = {'C': {0.001, 0.01, 0.1, 1, 10, 100}}
 - Allow the model to identify the optimal complexity level without restricting the search space too narrowly.
- Class_weight='balanced' is applied to mitigate the imbalance issues.
- GridSearchCV_log: {'C': 0.01} has stronger regularisation and not lead to underfitting and to learn in a more stable manner.

c. Report the variables that are included in the regression model.

```
coef = model.coef_[0]

coef = coef[:20]
for i in range(len(coef)):
    print(feature_names[i], ':', coef[i])

VehYear : -0.3804823517004311
VehOdo : 0.14415402803274585
MMRCurrentAuctionAveragePrice : -0.095570436780107
MMRCurrentRetailAveragePrice : -0.038272248876095455
Auction_ADESA : 0.07183110288718436
Auction_MANHEIM : -0.08946405705092635
Auction_OTHER : 0.03229484736480808
Transmission_AUTO : 0.03133792395509057
Transmission_MANUAL : -0.0318394074556094
Transmission_UNKNOWN : 0.00011396248691719884
Nationality_AMERICAN : -0.03239854940317287
Nationality_OTHER : 0.02095312916143133
Nationality_OTHER ASIAN : 0.041856578538101895
Nationality_TOP LINE ASIAN : -0.009323679205518009
Nationality_UNKNOWN : 0.00012395370954277815
```

d. Report the top-5 important variables (in order) in the model.

```

coef = model.coef_[0]

indices = np.argsort(np.absolute(coef))
indices = np.flip(indices, axis=0)

indices = indices[:20]
for i in indices:
    print(feature_names[i], ':', coef[i])

VehYear : -0.3804823517004311
Vehodo : 0.14415402883274585
MMRCurrentAuctionAveragePrice : -0.095570436780107
Auction_MANHEIM : -0.08946405705892635
Auction_ADESA : 0.07183110288718436
Nationality_OTHER ASIAN : 0.041856578538101895
MMRCurrentRetailAveragePrice : -0.038272248876095455
Nationality_AMERICAN : -0.03239854948317287
Auction_OTHER : 0.03229484736480888
Transmission_AUTO : 0.03133792395509057
Nationality_OTHER : 0.02095312916143133
Transmission_MANUAL : 0.0318394874556094
Nationality_TOP LINE ASIAN : -0.009323679205518009
Nationality_UNKNOWN : 0.00012395370954277815
Transmission_UNKNOWN : 0.00011396248691719884

```

1. VehYear : 0.380
2. VehOdo: 0.144
3. MMRCurrentAuction
AveragePrice: 0.096
4. Auction_MANHEIM: 0.0895
5. Auction_ADESA: 0.0718

e. What is the classification accuracy on training and test datasets?

Logistic		Regression_normal		GridSearch_cv		GridSearch_cv_log	
Training accuracy:	0.8704921985327042	Train accuracy:	0.6144842985568147	Train accuracy:	0.618916465469638	Test accuracy:	0.6204291569557181
Test accuracy:	0.8705296150446034	precision	recall	precision	recall	f1-score	support
		0	1	0	1	0.91	0.62
		0.87	0.00	0.61	0.19	0.74	10832
		1.00	0.00	0.61	0.61	0.74	1611
accuracy				accuracy		0.61	12443
macro avg	0.44	0.50	0.47	macro avg	0.55	0.61	12443
weighted avg	0.76	0.87	0.81	weighted avg	0.82	0.61	12443
				{'C': 1}			{'C': 0.01}

- Although the normal logistic regression shows a high accuracy (0.870), recall and f1-score are 0.00, which means that it completely failed to learn due to data imbalance.
- After applying GridSearch without log transformation, the accuracy decreased to 0.615, but recall and f1-score increased, which means it started to capture the minority class.
- Additionally, GridSearch with log transformation has slightly higher accuracy (0.620) so that the log transformation stabilised the feature distribution and allowed more effective regularisation.

f. Report any sign of overfitting.

- There is no accurate overfitting sign in the results of GridSearch and GridSearch with log transformation since the differences between train and test accuracy are very small (< 0.01).

3. Build another regression model on the reduced variables set. To minimise variables, either perform dimensionality reduction with Recursive Feature Elimination or select a subset of inputs found significant by the decision tree (use the best decision tree model under Task 2). Tune the model with GridSearchCV to find the best parameter setting. Answer the following:

a. Was dimensionality reduction helpful in identifying a good feature set for building the accurate model? Report the feature selection method used.

- Recursive Feature Elimination with Logistic Regression: Feature set is reduced from 15 to 2.
Original feature set 15
Number of features after elimination 2

- SelectedFromModel with Decision Tree: Feature set remains 3 from 15.

```
from sklearn.feature_selection import SelectFromModel

selectmodel = SelectFromModel(dt_best, prefit=True)
X_train_sel_model = selectmodel.transform(X_train)
X_test_sel_model = selectmodel.transform(X_test)
print(X_train_sel_model.shape)

(29033, 3)
```

→ By reducing the feature set, the model is simplified and improves the interpretation possibility.

b. Report the variables that are included in the regression model.

- Recursive Feature Elimination with Logistic Regression: Feature set is reduced from 15 to 2.

```
selected_features = feature_names[rfe.support_]
print("Selected features: ", selected_features)

Selected features: Index(['VehYear', 'Auction_MANHEIM'], dtype='object')
```

- SelectedFromModel with Decision Tree: Feature set remains 3 from 15.

```
VehYear : 0.7598540857360376
MMRCURRENTAuctionAveragePrice : 0.1144055334609266
VehOdo : 0.07466742514630023
Auction_ADESA : 0.03162454160340716
Auction_MANHEIM : 0.019448414053328502
Nationality_TOP LINE ASIAN : 0.0
Nationality_UNKNOWN : 0.0
Transmission_MANUAL : 0.0
Transmission_UNKNOWN : 0.0
Nationality_AMERICAN : 0.0
Nationality_OTHER : 0.0
Nationality_OTHER ASIAN : 0.0
Auction_OTHER : 0.0
Transmission_AUTO : 0.0
MMRCURRENTRetailAveragePrice : 0.0
Number of leaves: 8
```

c. What is the classification accuracy on the training and test datasets?


rfe_cv				rfe_cv_log				SelectFromModel_DT			
Train accuracy: 0.6316949677952675				Original feature set 15				Train accuracy: 0.68820445699721			
Test accuracy: 0.6333681588041469				Number of features after elimination 2				Test accuracy: 0.61440167162255991			
Precision	recall	f1-score	support	Train accuracy: 0.6316949677952675				precision	recall	f1-score	support
0	0.91	0.64	0.75	10832				0	0.91	0.62	0.74
1	0.19	0.57	0.29	1611				1	0.19	0.60	0.29
accuracy		0.63		12443				accuracy		0.61	
macro avg	0.55	0.61	0.52	12443				macro avg	0.55	0.61	0.51
weighted avg	0.82	0.63	0.69	12443				weighted avg	0.82	0.61	0.68
{'C': 0.001}				Best parameters: {'C': 0.001}				{'C': 0.001}			

- There are no differences of train (0.632) and test (0.633) accuracy between rfe_cv and rfe_cv_log.
 - Applying DT model leads to lower accuracy than RFECV.
- RFE models have better results than dt model since RFE chose two variables it decreases noise by removing unnecessary features.

d. Report any sign of overfitting.

- The accuracy of the three results is low (around 61-63%), and both train and test accuracy values are similar, which means there is no overfitting

4. Produce the ROC curve for both regression models. Using the best regression model, can you provide the characteristics of cars most likely to be ‘kicks’? If it is hard to comprehend, discuss why.



Task 4. Predictive Modelling Using Neural Networks

1. Describe what additional processing was required on this dataset to be used for neural network modelling.

```
from sklearn.preprocessing import StandardScaler
random_state = 10


scaler = StandardScaler()
X_train = scaler.fit_transform(X_train, y_train)
X_test = scaler.transform(X_test)
```

- Convert categorical data through dummies
- Normalise the numeric data to a 0-1 range by using 'scaler'

2. Build a Neural Network model tuned with GridSearchCV. Answer the following:

- a. Explain the parameters in building this model, e.g., network architecture, iterations, activation function, etc. Explain your choice of hyperparameters to search, and the chosen search range(s)?

<pre>from sklearn.neural_network import MLPClassifier from sklearn.metrics import classification_report model = MLPClassifier(random_state=random_state) model.fit(X_train_res, y_train_res) MLPClassifier(random_state=10)</pre>	<ul style="list-style-type: none"> Import MLP Classifier(Default neural network in sklearn) Solver hyperparameter is set by default to 'adam' and also activation is set by 'relu' automatically.
<pre>model = MLPClassifier(random_state=random_state) model.fit(X_train_res, y_train_res) C:\Users\ch93\anaconda\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:891: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet. warnings.warn(MLPClassifier(random_state=10) model = MLPClassifier(max_iter=500, random_state=random_state) model.fit(X_train_res, y_train_res) MLPClassifier(max_iter=500, random_state=10)</pre>	<ul style="list-style-type: none"> Due to warning message “convergence is not reached”, changed ‘max_iter’ hyperparameter default 200 to 500.



- We attempted to tune the model with the lower number of neurons in the hidden layer (less complex model).
- According to the cross-validation results, we set hidden layer size maximum 64. Beyond this point, the gap between test and train increased noticeably, indicating a higher risk of overfitting.
- With using 'alpha' hyperparameter, reduce overfitting by constraining the size of weights.

b. What is the classification accuracy of the training and test datasets?

```

print("Train accuracy:", model.score(X_train, y_train))
print("Test accuracy:", model.score(X_test, y_test))
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))

Train accuracy: 0.870644163538043
Test accuracy: 0.8705296150466934
      precision    recall   f1-score   support
          0       0.87     1.00     0.93    10832
          1       0.00     0.00     0.00    1611

accuracy                           0.87
macro avg       0.44     0.50     0.47    12443
weighted avg    0.76     0.87     0.81    12443

C:\Users\ch91\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is 1-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
C:\Users\ch91\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is 1-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
C:\Users\ch91\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is 1-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))

from imblearn.over_sampling import SMOTE
smote = SMOTE (random_state = 10)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)
# For fixing low precision, recall, f1-score
X_train_res = scaler.fit_transform(X_train_res) # scaling X_train_res

```

```

Train accuracy: 0.6623273849562774
Test accuracy: 0.5528409547536768
      precision    recall   f1-score   support
          0       0.91     0.54     0.68    10832
          1       0.17     0.63     0.27    1611

accuracy                           0.55
macro avg       0.54     0.58     0.47    12443
weighted avg    0.81     0.55     0.63    12443

```

- Warning message**
Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use zero_division parameter to control this behavior. `_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))`
- This dataset has a very low ratio of IsBadBuy=1 (defective purchase), causing a problem that the model can hardly predict minority classes with basic learning.
- Although Accuracy itself is high, the imbalance between data is quite severe, and to solve this problem, **SMOTE (Synthetic Minority Oversampling Technique)** was applied in the learning stage.
- SMOTE balances between classes while preserving data distribution by generating new synthetic samples among existing samples, rather than simply replicating minority class data.
- Through this process, the model was able to learn both classes evenly, and in particular, the Recall and F1-score for detecting IsBadBuy=1 were improved.

```

print("Train accuracy:", cv_1.score(X_train_res, y_train_res))
print("Test accuracy:", cv_1.score(X_test, y_test))

y_pred = cv_1.predict(X_test)
print(classification_report(y_test, y_pred))
print(cv_1.best_params_)

Train accuracy: 0.6791239662881335
Test accuracy: 0.5979265450454071
      precision    recall   f1-score   support
          0       0.90     0.60     0.72     10832
          1       0.17     0.55     0.26     1611

           accuracy          0.60     12443
    macro avg       0.54     0.58     0.49     12443
weighted avg       0.81     0.60     0.66     12443

{'hidden_layer_sizes': (240,)}

# new parameters
params = {'hidden_layer_sizes': [(12,), (16,), (24,), (32,), (40,), (48,), (56,), (64,)]}
cv_2 = GridSearchCV(param_grid=params,
                     estimator=MLPClassifier(random_state=random_state), cv=10, n_jobs=-1)
cv_2.fit(X_train_res, y_train_res)

print("Train accuracy:", cv_2.score(X_train_res, y_train_res))
print("Test accuracy:", cv_2.score(X_test, y_test))

y_pred = cv_2.predict(X_test)
print(classification_report(y_test, y_pred))
print(cv_2.best_params_)

Train accuracy: 0.6450362046452736
Test accuracy: 0.5968014144498915
      precision    recall   f1-score   support
          0       0.91     0.60     0.72     10832
          1       0.18     0.58     0.27     1611

           accuracy          0.60     12443
    macro avg       0.54     0.59     0.50     12443
weighted avg       0.81     0.60     0.66     12443

{'hidden_layer_sizes': (64,)}

params = {'hidden_layer_sizes': [(12,), (16,), (24,), (32,), (40,), (48,), (56,), (64,)],
          'alpha': [0.01, 0.001, 0.0001, 0.00001]}
cv_3 = GridSearchCV(param_grid=params,
                     estimator=MLPClassifier(random_state=random_state), cv=10, n_jobs=-1)
cv_3.fit(X_train_res, y_train_res)

print("Train accuracy:", cv_3.score(X_train_res, y_train_res))
print("Test accuracy:", cv_3.score(X_test, y_test))

y_pred = cv_3.predict(X_test)
print(classification_report(y_test, y_pred))
print(cv_3.best_params_)

Train accuracy: 0.6450362046452736
Test accuracy: 0.5968014144498915
      precision    recall   f1-score   support
          0       0.91     0.60     0.72     10832
          1       0.18     0.58     0.27     1611

           accuracy          0.60     12443
    macro avg       0.54     0.59     0.50     12443
weighted avg       0.81     0.60     0.66     12443

{'alpha': 0.0001, 'hidden_layer_sizes': (64,)}

```

CV1

- Train accuracy ≈ 0.68, Test accuracy ≈ 0.60
- The output of this GridSearchCV returns 240 neurons as the optional number of neurons in the hidden layer.
- But too many hidden nodes can lead overfitting

CV2

- Train accuracy ≈ 0.65, Test accuracy ≈ 0.60
- In previous models, there were cases where the less complex model produced better results. Therefore, as a result of trying a smaller number of hidden layers even in NN, the Accuracy difference decreased, the recall value increased, and a lower number of hidden layers was returned.
- The output of this GridSearchCV returns 64 neurons as the optional number of neurons in the hidden layer.

CV3

- Train accuracy ≈ 0.65, Test accuracy ≈ 0.60
- Tuned the second hyperparameter, 'alpha' for strength the L2 regularization term.
- The default value of alpha is 0.0001, thus we tried alpha options around this value.
- There's not a remarkable difference with CV2.

c. Did the training process converge and result in the best model?

- When code is executed with the default MLP Classifier (max_iter=200), a warning message is generated, increased to max_iter=500, and data is collected within 500 times without an error message.

d. Do you see any sign of over-fitting?

```
Train accuracy: 0.6450362046452736
Test accuracy: 0.5968014144498915
      precision    recall  f1-score   support
0       0.91     0.60     0.72     10832
1       0.18     0.58     0.27     1611

accuracy                           0.60    12443
macro avg                           0.54    12443
weighted avg                          0.81    12443

{'alpha': 0.0001, 'hidden_layer_sizes': (64,)}
```

- Although it is difficult to regard it as overfitting with a difference of about 5% in accumulation, it can be interpreted that the accumulation itself is low, and the amount of learning of the model tends to be insufficient.

3. Build another Neural Network model with the reduced feature set. Perform dimensionality reduction by selecting variables with a decision tree (use the best decision tree model under Task 2). Tune the model with GridSearchCV to find the best parameter settings. Answer the following:

a. Did feature selection favour the outcome? Report the changes in the network architecture. What inputs are being used as the network input?

```
import pickle
from sklearn.feature_selection import SelectFromModel

with open('decision_tree_model.pickle', 'rb') as f:
    dt_best, roc_index_dt_cv, fpr_dt_cv, tpr_dt_cv = pickle.load(f)

selectmodel = SelectFromModel(dt_best, prefit=True)
X_train_sel_model = selectmodel.transform(X_train_res)
X_test_sel_model = selectmodel.transform(X_test)
print(X_train_sel_model.shape)
```

Select only the significant variables selected in the DT model

- VehYear : 0.7661573891371191
- MMRCurrentAuctionAveragePrice : 0.11149221549044215
- VehOdo : 0.0731100508607268
- Auction_ADESA : 0.029442055720801594
- Auction_MANHEIM : 0.01979828879091031

```
print(X_train_sel_model.shape)
(50546, 3)
```

- Simplified input data by selecting the feature set.
- Features reduced from 15 to 3

b. What is the classification accuracy on the training and test datasets?

```
params = {'hidden_layer_sizes': [(12,), (16,), (24,), (32,), (40,), (48,), (56,), (64,)],
          'alpha': [0.01, 0.001, 0.0001, 0.00001]}
cv_sel_model = GridSearchCV(param_grid=param_params,
                           estimator=MLPClassifier(random_state=random_state), cv=10, n_jobs=-1)
cv_sel_model.fit(X_train_sel_model, y_train_res)

print("Train accuracy:", cv_sel_model.score(X_train_sel_model, y_train_res))
print("Test accuracy:", cv_sel_model.score(X_test_sel_model, y_test))

y_pred = cv_sel_model.predict(X_test_sel_model)
print(classification_report(y_test, y_pred))
print(cv_sel_model.best_params_)
```

Train accuracy: 0.6176354212004906
Test accuracy: 0.53612472876316
 precision recall f1-score support
0 0.92 0.51 0.66 10832
1 0.18 0.71 0.28 1611

accuracy 0.54 12443
macro avg 0.55 12443
weighted avg 0.82 12443

{'alpha': 1e-05, 'hidden_layer_sizes': (48,)}

- Train accuracy ≈ 0.62, Test accuracy ≈ 0.54
- The DT-based reduced feature NN recorded a training accuracy of about 0.61 and a test accuracy of about 0.56.
- Although the overall accuracy is slightly lower than that of the entire feature model, it can be interpreted that the recall of the minority class (IsBadBuy=1) is improved by about 2%, so that a defective purchase vehicle can be detected slightly better


c. How many iterations are now needed to train this network?

- Epoch setting remained at max_iter=500, converging stably within 500 times

d. Do you see any sign of over-fitting? Did the training process converge and result in the best model?

<pre> Train accuracy: 0.6176354212004906 Test accuracy: 0.53612472876316 precision recall f1-score support 0 0.92 0.51 0.66 10832 1 0.18 0.71 0.28 1611 accuracy 0.54 12443 macro avg 0.55 0.61 0.47 12443 weighted avg 0.82 0.54 0.61 12443 {'alpha': 1e-05, 'hidden_layer_sizes': (48,)}</pre>	<ul style="list-style-type: none"> Although this result has a slightly larger difference in Train/Test accuracy than the previous model, it is still not possible to determine the overfitting at 8%. In addition, both models converged stably, and although this model is slightly less accurate, our target detection rate of defective purchases itself is a little higher, so this model can be seen as a more optimal model.
---	--

4. Produce the ROC curve for both NNs. Using the best NN model, can you provide the characteristics of cars most likely to be 'kicks'? If it is hard to comprehend, discuss why.




Task 5. Final remarks: The decision making

1. Finally, based on all models and analysis, is there a model you will use in decision-making i.e. the best-performing model? Justify your choice. Draw an ROC chart and an Accuracy Table to support your findings.

<ROC chart>

ROC index on test for decision tree: 0.6386763716857975
 ROC index on test for linear regression: 0.645443656388112
 ROC index on test for NN with feature selection (model selection) and gridsearch: 0.6471949677576704

**<Accuracy Table>**

Model	Training Accuracy	Test Accuracy	Precision (class 1)	Recall (class 1)	F1-score (class 1)	ROC
Decision Tree	0.627	0.623	0.19	0.57	0.28	0.639
Linear Regression	0.608	0.614	0.19	0.60	0.29	0.645
Neural Network	0.617	0.536	0.18	0.71	0.28	0.647

Best-performing model: Neural Network (NN) with cv_sel_model

- This model has the highest ROC index (0.647), Recall (0.71).
- In the neural network model, the overall accuracy is relatively low, but the recall for IsBadBuy=1 is comparatively high. In other words, it means that the ability to detect kick vehicles is high, and it is a number that matches our goal better. This result will be more helpful for dealers to predict defects.
- According to the prediction accuracy table, this model can predict class 1 ("kick") better than other models.
- A decision tree has vulnerability related to overfitting and data imbalance. This model cannot be selected for stability because our dataset has these issues.
- Logistic regression cannot capture complex nonlinear relationships, and it can lead to lower accuracy and precision-recall imbalance compared to neural networks.

2. Based on this analysis, can you summarise the positives and negatives of each predictive modelling method?

Model	Pros	Cons
Decision Tree	<ul style="list-style-type: none"> ▪ Can explain reasons about relationship with variables (in the node) ▪ Low necessity of pre-processing (scaling) 	<ul style="list-style-type: none"> ▪ Vulnerability of overfitting and data imbalance: our dataset has severe data imbalance, which can lead to overfitting
Logistic Regression	<ul style="list-style-type: none"> ▪ Simple and highly interpretable ▪ Easily identify key variables ▪ Logarithmic transformation: Reduce skewness 	<ul style="list-style-type: none"> ▪ Precision-recall imbalance: recall > precision ▪ Model Complexity: Linear regression
Neural Network	<ul style="list-style-type: none"> ▪ Neural Network can handle the non-linear data even with missing values. 	<ul style="list-style-type: none"> ▪ Takes a long time for training and is sensitive to the number of nodes and other parameters. ▪ Also, cannot interpret exactly decipher the process inside (such as black-box).